

# Machine Intelligence:: Deep Learning

*Beate Sick, Lilach Goren, Pascal Bühler*



Institut für Datenanalyse und Prozessdesign  
Zürcher Hochschule für Angewandte Wissenschaften

# Outline of the DL Module (tentative)

The course is split in 8 sessions, each 4 lectures long. Topics might be adapted during the course

| Day | Date       | Time        | Topic   |
|-----|------------|-------------|---|
| 1   | 15.04.2025 | 09:00-12:30 | Introduction to Deep Learning & Keras, first NNs  |
| -   | 21.04.2025 | -           | FRÜHLINGS-FERIEN                                  |
| -   | 28.04.2025 | -           | FRÜHLINGS-FERIEN                                  |
| 2   | 06.05.2025 | 09:00-12:30 | Loss, Optimization, Regression, Classification    |
| 3   | 13.05.2025 | 09:00-12:30 | Computer vision, CNN-architecture                 |
| 4   | 20.05.2025 | 09:00-12:30 | DL in practice, pretrained (foundation) models    |
| 5   | 27.05.2025 | 09:00-12:30 | Model evaluation, baselines, xAI, troubleshooting |
| 6   | 03.06.2025 | 09:00-12:30 | Generative Models, Transformer-architecture       |
| 7   | 10.06.2025 | 09:00-12:30 | Vision Transformer                                |
| 8   | 17.06.2025 | 09:00-12:30 | Projects, deep Ensembling                         |

## Look back on homework: MNIST with CNN



Use a CNN to classify MNIST data to 10 possible classes

[06\\_cnn\\_mnist\\_shuffled\\_keras\\_torch.ipynb](#) <- first part with the unshuffled cnn

| Model | Loss | Accuracy | # Params |
|-------|------|----------|----------|
| FcNN  | 0.1  | 0.97     | 84060    |
| CNN   | 0.05 | 0.98     | 35962    |


What is a good  
CNN architecture?

# CNN breakthrough in 2012: Imagenet challenge

1000 classes  
1 Mio samples



...




Human: 5% misclassification

Only one non-CNN approach in 2013

GoogLeNet 6.7%

A. Krizhevsky  
first CNN in 2012  
**Und es hat zoom gemacht**



# Review of ImageNet winning CNN architectures




# The trend in modern CNN architectures goes to small filters

Why do modern architectures use very small filters?  
Determine the receptive field in the following situation:


1) Suppose we have **one**  
**7x7 conv layers** (stride 1)  
**49 weights**



2) Suppose we **stack two**  
**3x3 conv layers** (stride 1)



3) Suppose we **stack three**  
**3x3 conv layers** (stride 1)  
**3\*9=27 weights**



**We need less weights for the same receptive field when stacking small filters!**

# "Oxford Net" or "VGG Net" 2014 2<sup>nd</sup> place

- 2<sup>nd</sup> place in the imageNet challenge
- More **traditional**, easier to train
- Small pooling
- **Stacked 3x3 convolutions before maxpooling**  
-> **large receptive field**
- no strides (stride 1)
- ReLU after conv. and FC (batchnorm was not introduced)
- Pre-trained model is available (see exercise)

<http://arxiv.org/abs/1409.1556>



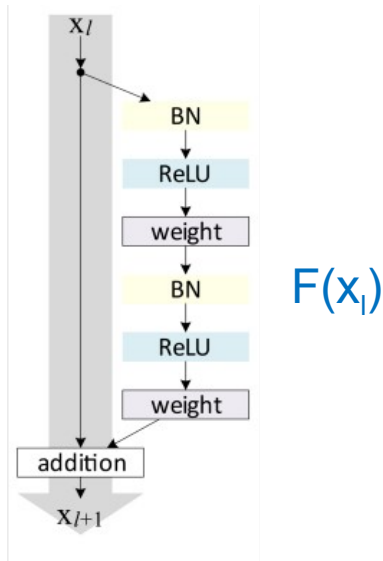


## "ResNet" from Microsoft 2015 winner of imageNet

152  
layers

## ResNet basic design (VGG-style)

- add shortcut connections every two
- all 3x3 conv (almost)

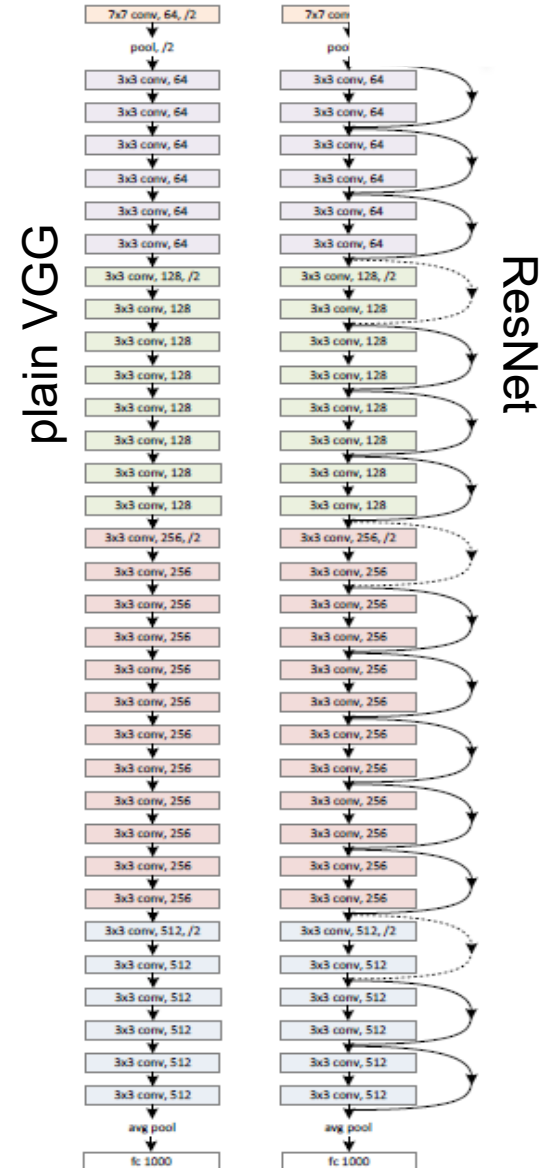


$$H(x_i) = x_{i+1} = x_i + F(x_i)$$

$F(x)$  is called “residual” since it only learns the “delta” which is needed to add to  $x$  to get  $H(x)$

152 layers:  
Why does this train at all?

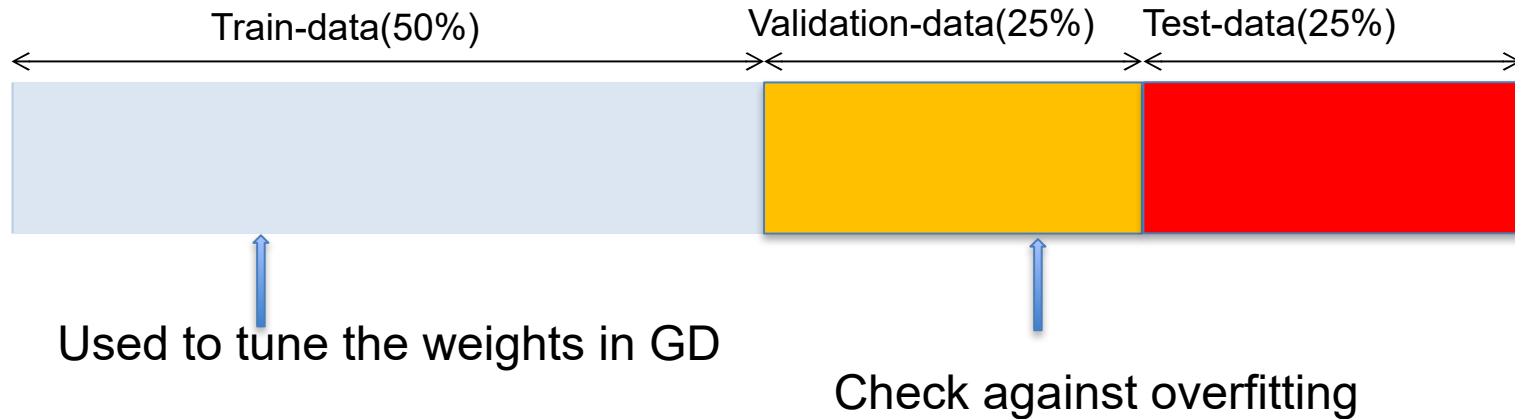
This deep architecture could still be trained, since the gradients can skip layers which diminish the gradient!



# Tricks of the trade

- Early stopping
- Input Standardization
- Batch Norm Layer
- Dropout
- Data augmentation

# Best practice: Split in Train, Validation, and Test Set



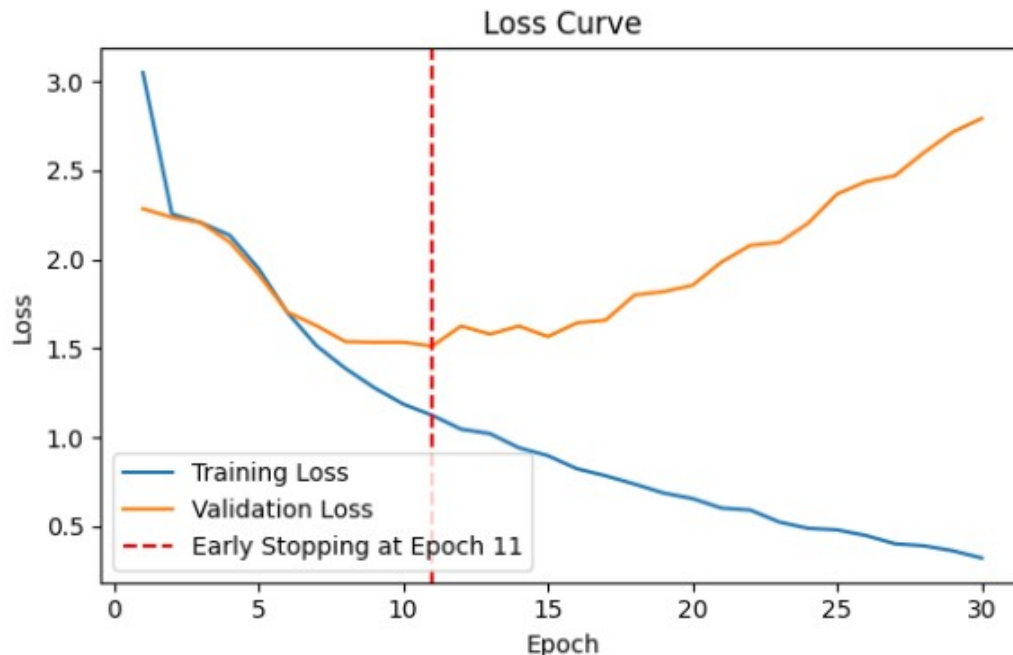
Best practice: Lock an extra **test data set** away, and use it only at the very end, to evaluate the chosen model, that performed best on your validation set.

Reason: **When trying many models, you probably overfit on the validation set.**

Determine performance metrics, such as MSE, to evaluate the predictions **on new validation or test data**

# Loss curve and early stopping

Very common check: Plot loss in train and validation data vs epoch of training.



Training completed all epochs.  
Best model weights were restored from epoch 11

```
## Early Stopping
early_stopping = EarlyStopping(
    monitor='val_loss',
    patience=10,
    verbose=1,
    restore_best_weights=False
)
```



- If **training loss** does not go down to zero: model is not flexible enough
- Use weights @minimum of validation before overfitting
- Early stopping: stop to training if validation loss does not improve anymore

[07 early stopping and modelweights](#)



# Take-home messages from CIFAR10 CNN study

- DL does not need a lot of preprocessing, but working with standardized (small-valued) input data often helps.

Without normalizing  
the input to the CNN



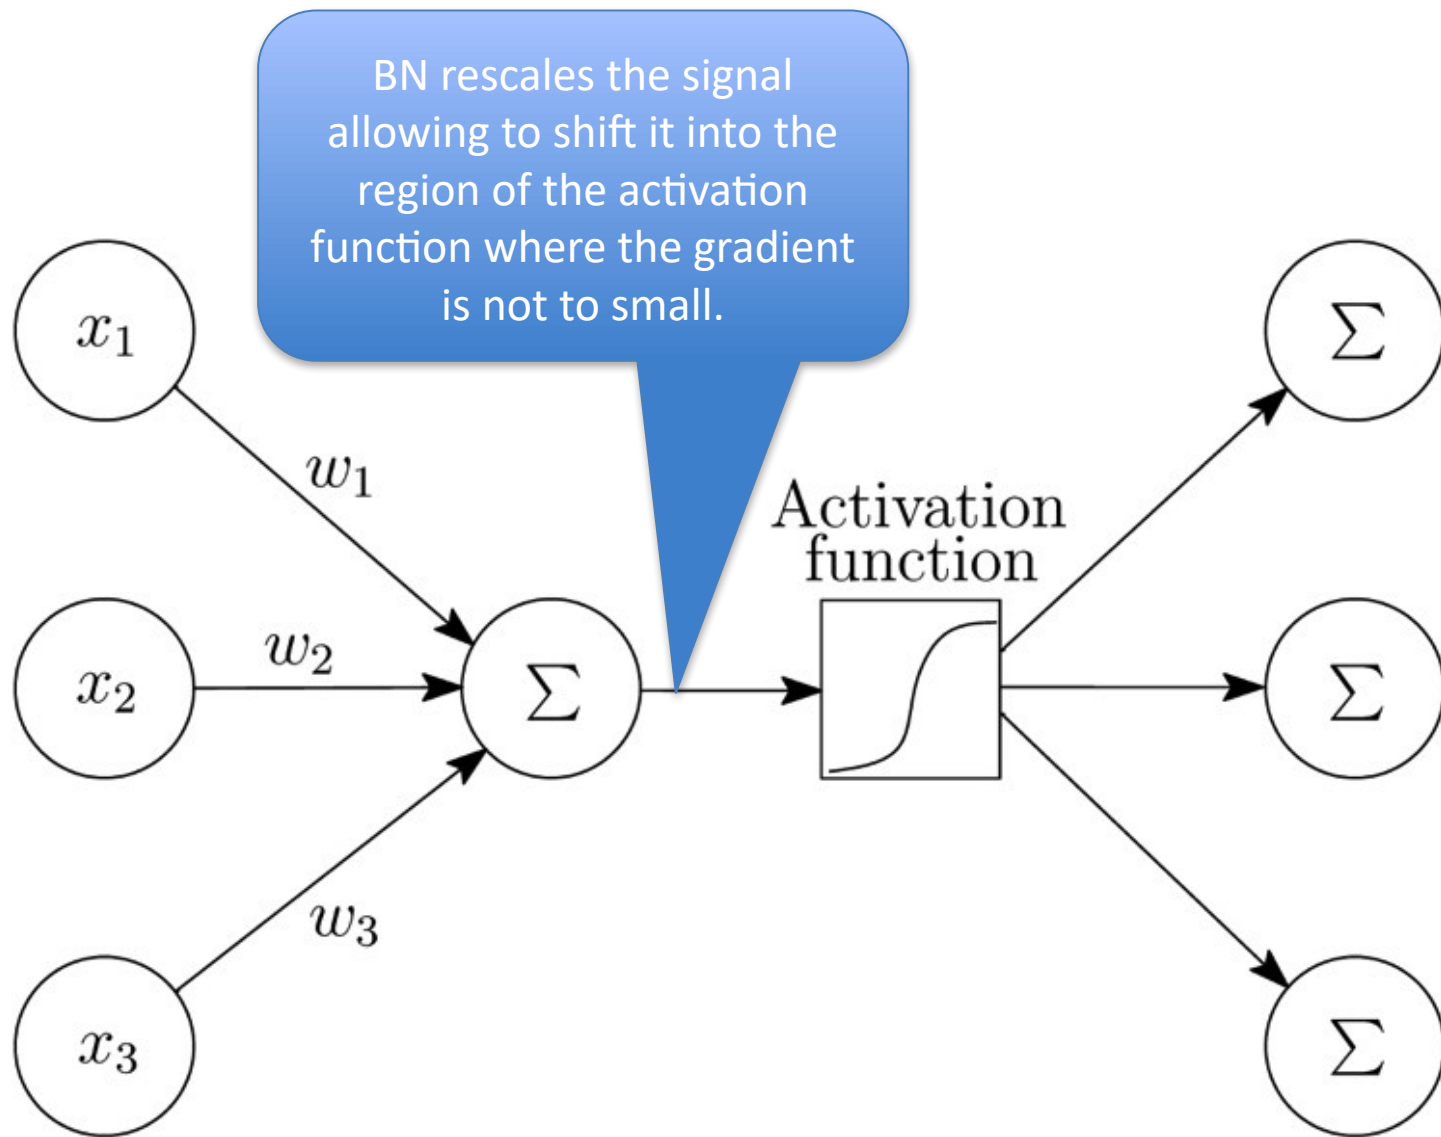
With normalizing  
(pixel-value/255)  
the input to the CNN




## Regularization to avoid overfitting

- Batchnorm layer
- Dropout layer

# What is the idea of Batch-Normalization (BN)



# Dropout helps to fight overfitting




Using **dropout** during training implies:

- In each training step only weights to not-dropped units are updated  $\rightarrow$  we train a sparse sub-model NN
- For predictions with the trained NN we freeze the weights corresponding to averaging over the ensemble of trained models we should be able to “reduce noise”, “overfitting”
- JFI: To get same expected output in training (with dropout) and after training (test time - without dropout), the weights are multiplied after training by the dropout probability  $p=0.5$ .





# Dropout



Three NNs:

a) shows the full NN with all neurons (as used when NN is trained),  
b) and c) show two versions of a thinned NN where some neurons are dropped (as done during training with dropout). Dropping neurons is the same as setting all connections that start from these neurons to zero.

# Dropout fights overfitting in a CIFAR10 CNN



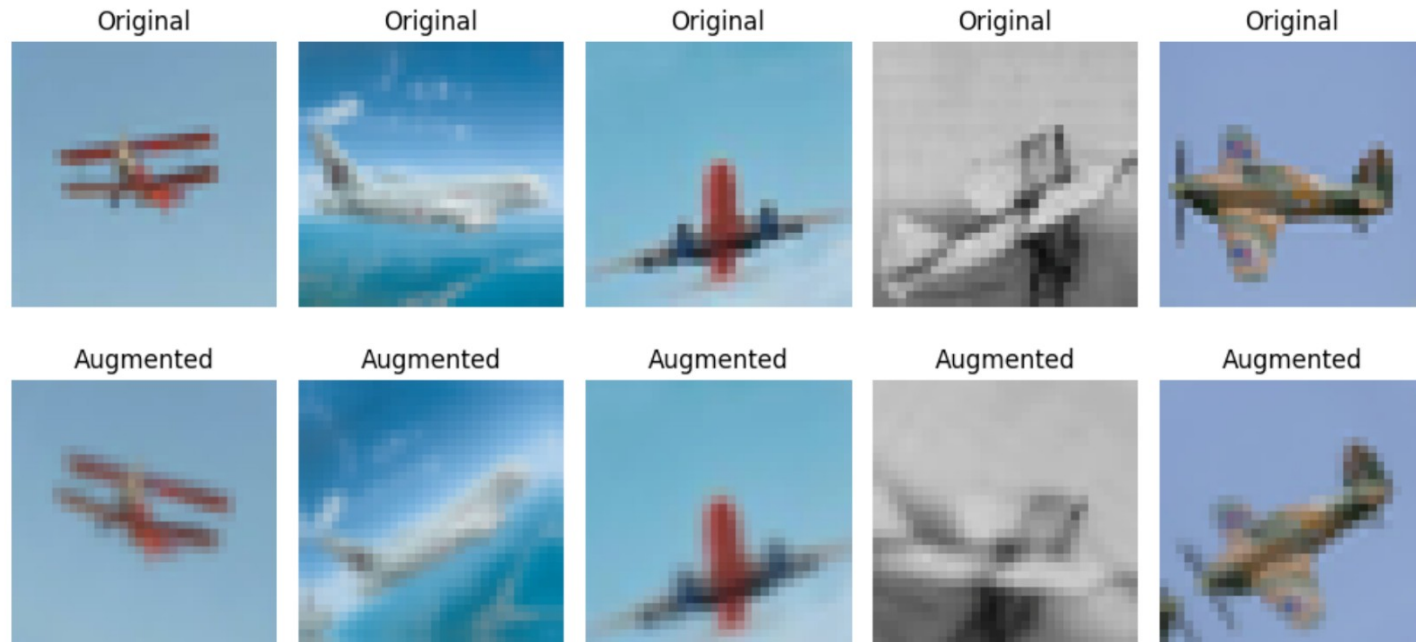
# What to do in case of limited data?

- Data augmentation
- Pretrained (foundation) models

# Fighting overfitting by Data augmentation:

During training random operations are done on the fly

- Rotate image within an angle range
- Flip image: left/right, up, down
- resize
- Take patches from images
- ....



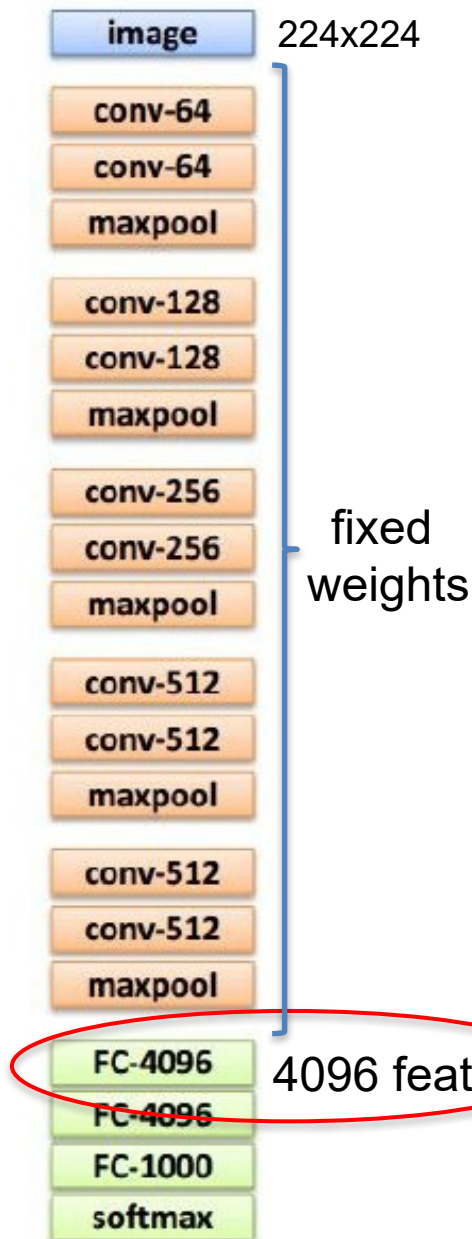
# In Code: Augmentation

```
1 # Define data augmentation pipeline
2 from keras.layers import RandomFlip, RandomRotation, RandomContrast, RandomZoom
3 data_augmentation = keras.Sequential([
4     #RandomFlip("horizontal"),      # Randomly flip images horizontally
5     RandomRotation(0.1),            # Randomly rotate images by 10%
6     RandomZoom(0.1),               # Randomly zoom in on images
7     #RandomContrast(0.1),          # Adjust contrast randomly
8 ], name='Augmentation')
9
```

```
1 # Functional API Model Definition
2 x = Input(shape=input_shape)
3 x = data_augmentation(x)
4 # First Convolutional Block
5 x = Convolution2D(8, kernel_size, padding='same')(x)
6 ...
```

[07\\_cifar10\\_tricks\\_keras\\_torch.ipynb](#)

# Use pre-trained CNNs for feature generation



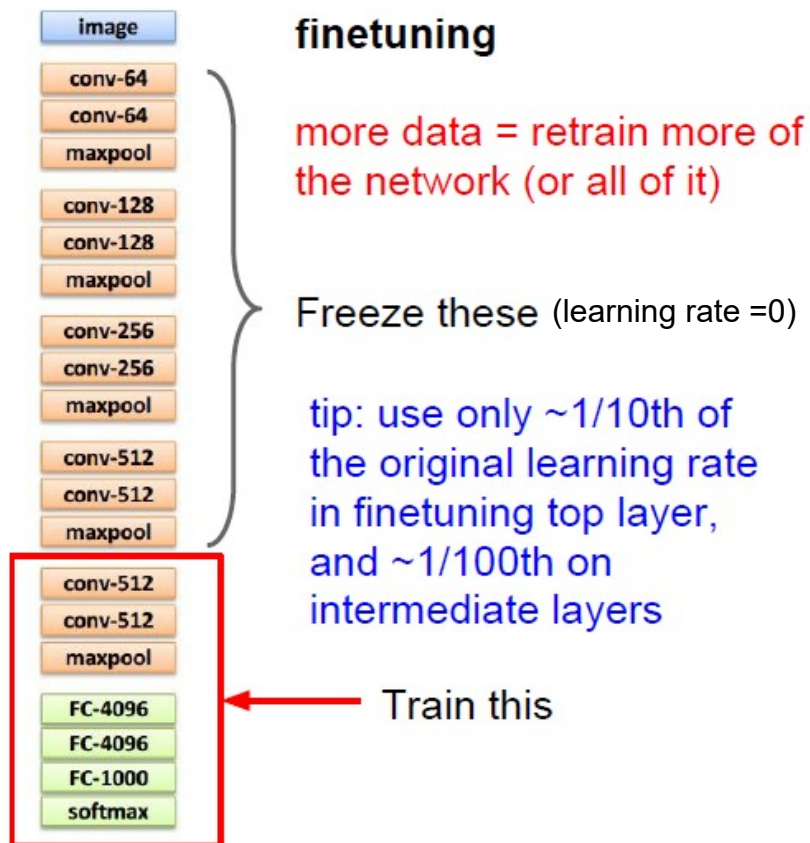
- Load a pre-trained CNN – e.g.g VGG16  
developed by Karen Simonyan and Andrew Zisserman in 2014
- Images are resized to required size
- Rescaling of the pixel values to “VGG range”
- Do a forward pass and **fetch features** that are used as CNN representations, dump these features into a file on disk
- Use these CNN features as input to a simple classifier – e.g. fc NN, RF, SVM ...  
(here it is easily possible to adapt to the new number of class labels)

Number features depends on input shape

Fetch this CNN feature vector for each image

# Transfer learning beyond using off-shelf CNN feature

e.g. medium data set (<1M images)



The strategy for fine-tuning depends on the size of the data set and the type of images:

|              | Similar task<br>(to imageNet challenge)   | Very different task<br>(to imageNet challenge)  |
|--------------|---|---|
| little data  | Extract CNN representation of one top fc layer and use these features to train an external classifier | You are in trouble - try to extract CNN representations from different stages and use them as input to new classifier |
| lots of data | Fine-tune a few layers including few convolutional layers   | Fine-tune a large number of layers  |


Where to get pretrained CNNs like VGG16 or ResNet: <https://keras.io/api/applications/>

Hint: first retrain only fully connected layer, only then add convolutional layers for fine-tuning.



# Performance of off-the-shelf CNN features when compared to tailored hand-crafted features

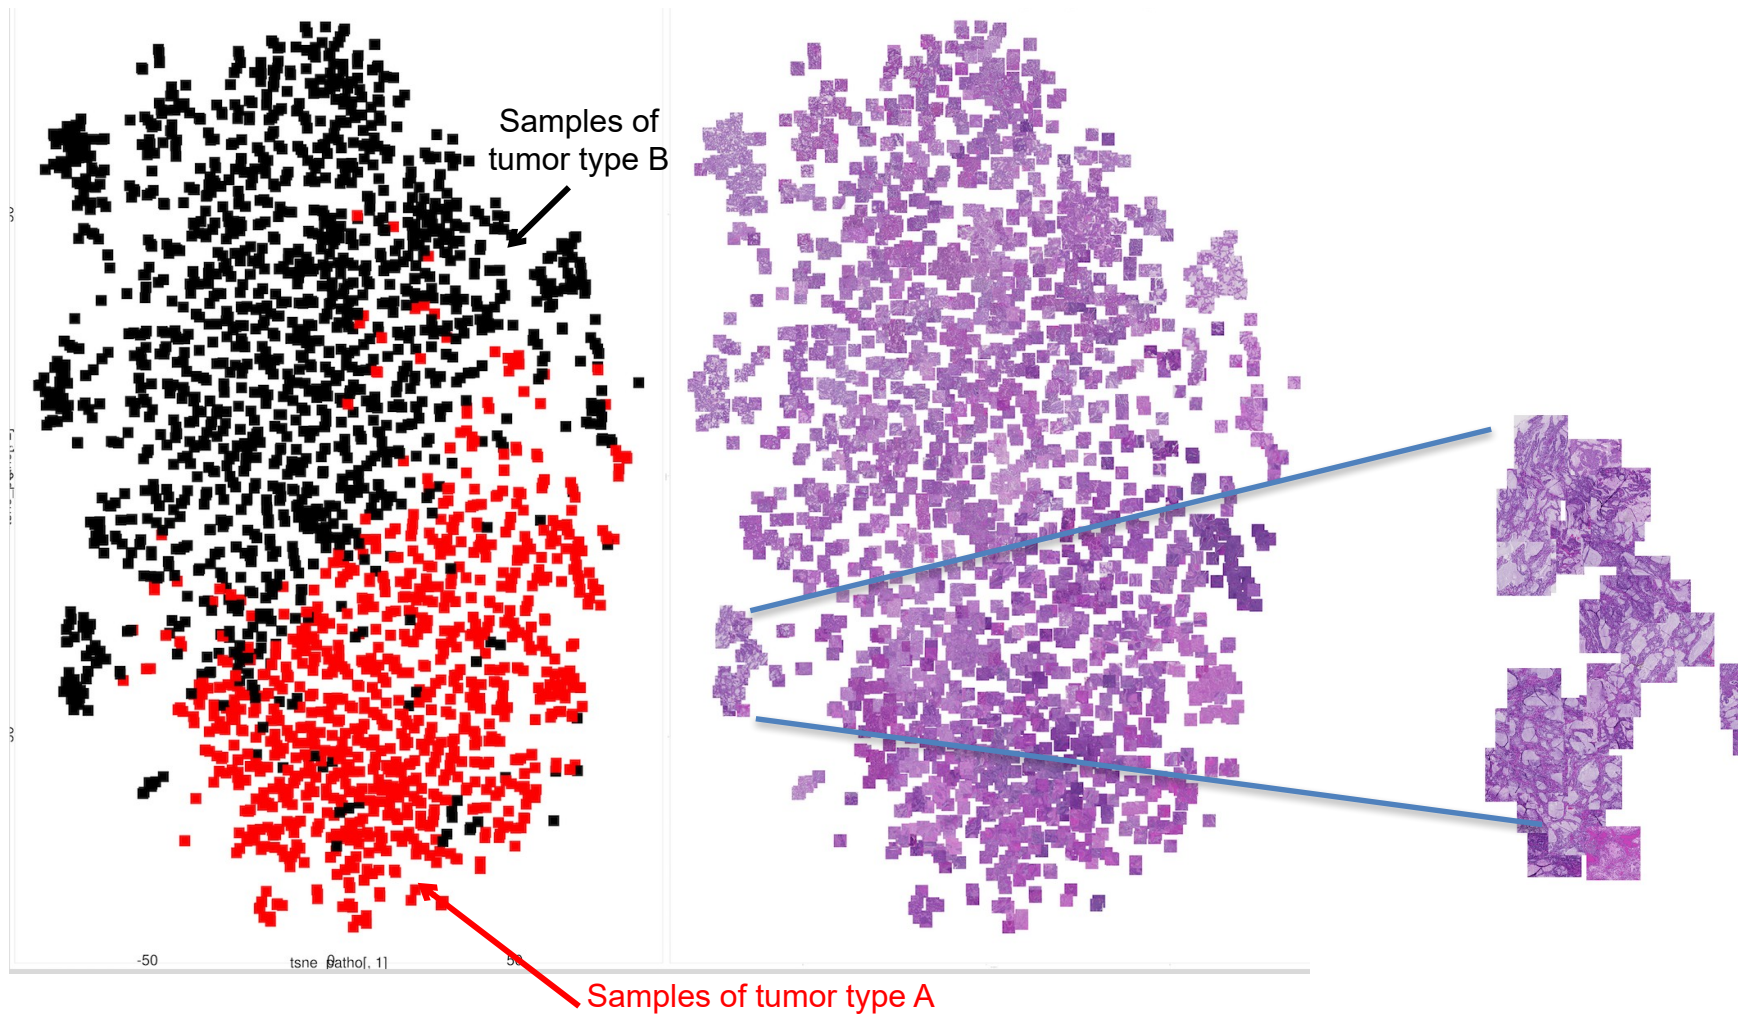
CNN's distributed and compositional features generalize well to new tasks



“Astonishingly, we report consistent superior results compared to the highly tuned state-of-the-art systems in all the visual classification tasks on various datasets.”




# Use features from a pretrained VGG as input to t-SNE




- Different tissue types cluster together in t-SNE: we could use knn as classifier
- VGG features even work on images that are far away from the 1000 imageNet classes

# Case Study: CIFAR10 Dataset

# Cifar10 Data



10 Classes, with 5'000 images for training



Some examples:

|                      |      |       |
|----------------------|------|-------|
| VGG-19 with GradInit | 2021 | 94.71 |
| ResNet-18            | 2022 | 95.55 |


# Dataset Splitting in Machine Learning Competitions

- Training Set
  - The set you get for training, you can do what you want with it.
- Test Set:
  - Final evaluation set, never seen by the model during training (e.g., 10,000 images in CIFAR-10).

Note: CIFAR-10 and many other datasets do NOT provide a separate validation set.

- Typical Workflow
  - Hyperparameter Tuning
    - **Create validation set** from training set
      - Helps in tuning hyperparameters and preventing overfitting.
      - Common practice: split 80% training / 20% validation (e.g., 40,000 train / 10,000 validation).
  - Final Training: After hyperparameter tuning, models are often retrained on the entire training set
  - Evaluation on test set
    - Sometimes test set is secret and you to upload you predictions (private test set)

```
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
```



# Exercise: Transfer learning with CIFAR10 data



# Take home: DL with images in practice

- Use NN architectures inspired by challenge winning NNs
- Use tricks of the trade when building a CNN
  - Normalization of input data
  - Batchnorm
  - Dropout
  - Data Augmentation
  - Early stopping (patience)
- In case of few data
  - Do augmentation during training of a CNN
  - Use shallow learner (e.g. RF) based on image features extracted via pretrained CNNs
  - Fine-tune a pretrained CNN on few data (transfer learning)
- Use pretrained (foundation) models